

Using PGFPLOTS to make economic graphs in L^AT_EX

Arnav Bandekar*

August 5, 2021




Figure 0-1: Wage elasticities and total income.

*If you have any questions, inquiries, or even suggestions on better ways to make graphs, contact me at bandekar.arnav@gmail.com by email. Do not hesitate to reach out.

Contents

1	Introduction	3
1.1	New to LaTeX?	3
1.2	GitHub project repository	3
1.3	Packages	4
2	Outlining the axes	4
2.1	The axis environment	5
2.2	Axis line directions	6
2.3	Axis ticks, scale, and clipping	7
2.4	Axis labels	8
3	Plotting, colours, and labelling (ex. indifference map and budget constraint)	10
3.1	Plotting functions	10
3.2	Colours	15
3.3	Plotting dashed line segments	17
3.4	Plotting coordinate points	18
3.5	Labelling	19
4	Empirical curves (ex. market equilibrium)	22
4.1	Major and minor grid lines	24
4.2	Distancing labels, colouring label backgrounds	26
5	Shading, arrows, and graphs side by side (ex. supply increase, price-taking firm)	29
5.1	Opacity and transparency	31
5.2	Shading in boxes with colours	32
5.3	Straight and curly arrows	34
5.4	Arranging graphs side by side	39
5.5	Shading in the area under a curve	45
6	Axis dimension lines, line breaks in labels, and a legend (ex. excise tax)	48
6.1	Adding dimension lines and line breaks in labels	51
6.2	Adding a legend	52
7	Conclusion	59
8	Resources	60
9	Other examples	62

1 Introduction

Economics makes an abundant use of graphs to illustrate important concepts and phenomena. Yet the first thing one may notice when typing up a paper in the field is the difficulty of making these graphs using word processors. That is where LaTeX comes in. This guide will explain how we can use the pgfplots package to make elegant economic graphs in LaTeX.

1.1 New to LaTeX?

Some readers may be unfamiliar with LaTeX. It is a powerful typesetting language capable of typesetting text, mathematical formulas, tables, and many other elements that word processors such as Word and Google Docs are incapable of. Its capabilities are extended by packages. This guide will focus on making graphs with a package designed for plotting graphs—pgfplots—which itself relies on a package designed for making graphics—TikZ.

Along with LaTeX’s formidable ability and customizability comes with it a reputable learning curve. Do not despair. This guide was written under the assumption that readers are familiar with the very basics of LaTeX—but no more than that.¹ If you can create a text document, load in packages, and write an equation in LaTeX’s display math mode, you need not know anything else. Although if you are already familiar with pgfplots and TikZ, you may still find some of the techniques used here useful for the very specific features of economic graphs.

1.2 GitHub project repository

All of the finished graphs shown here are available on the GitHub repository for this guide. You can access the repository at <https://github.com/jackypacky/pgf-econ-graphs>. There, the TeX source codes and PDFs of all the graphs can be found.

¹In Section 8: Resources, there are resources you can use to set up and learn the basics of LaTeX.

1.3 Packages

We will be using pgfplots to draw economic graphs, so the pgfplots package should be called. This will call TikZ as well. Since we want to be able to shade in areas under curves, the fillbetween pgfplots library should be called. Since we want to be able to specify positions around coordinates and use a variety arrowheads in TikZ, the positioning and arrows.meta TikZ libraries should be called as well. Also, since we want to be able to use LaTeX's math mode, we need to call the amsmath package. In total, we have the following code snippet at the beginning of the document:

```
\usepackage{pgfplots}
\pgfplotsset{compat = newest}
\usetikzlibrary{positioning, arrows.meta}
\usepgfplotslibrary{fillbetween}
\usepackage{amsmath}
```

Out of personal preference, we will be using the Times New Roman 12 point font with a 1.25 line spread (which corresponds to a one-and-a-half line spacing). So we will be using this document preamble:

```
\documentclass[12pt]{article}
\linespread{1.25}
\usepackage{times}

\usepackage{pgfplots}
\pgfplotsset{compat = newest}
\usetikzlibrary{positioning, arrows.meta}
\usepgfplotslibrary{fillbetween}
\usepackage{amsmath}
```

2 Outlining the axes

To begin, we need to draw the axes on which our economic graphs will be drawn. Since pgfplots is a dependency of TikZ, we need to open and close the TikZ environment with `\begin{tikzpicture}` and `\end{tikzpicture}`. Within this environment, to construct the pgfplots environment, we insert

`\begin{axis}` and `\end{axis}`. Since, presumably, we want the economic graph to be in the center of the page, all of this should be in the center environment, with `\begin{center}` and `\end{center}` encompassing the code. So we have:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
% Insert parameters here
]
% Insert commands here
\end{axis}
\end{tikzpicture}
\end{center}
```

For the axis environment, parameters affect the whole graph (i.e. scaling, axis ticks, grid lines) while commands impose specific plots (i.e. functions, lines, coordinate points). One thing to note, syntactically, is that parameters are separated by commas—“,”—while commands are separated by semicolons—“;”.

2.1 The axis environment

Before we can begin plotting functions, we need to specify which region of the Cartesian plane we wish to display. To specify the domain, we use the parameters `xmin = 0` and `xmax = 10`, where 0 and 10 are the minimum and maximum bounds of the domain. Similarly, to specify the range, we use the parameters `ymin = 0` and `ymax = 10` to specify the minimum and maximum range bounds respectively. So far, we have this code for the basic outline of the graph:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
]
% Insert commands here
\end{axis}
\end{tikzpicture}
\end{center}
```

```
\end{tikzpicture}  
\end{center}
```

Running this code produces Figure 2-1.




Figure 2-1: The x and y-axes, ticked and boxed.

2.2 Axis line directions

We need to make a couple modifications to make this suitable for economic graphs. The default style of the graph frame is boxed, where black lines surround the graph on all sides. While this is good for line graphs and scatter plots—what pgfplots was built in mind for—economic graphs generally have an L-shaped frame. To change the frame style, we need to use the parameter `axis lines = style`, where the style can either be `left`, `right`, `center`, `box`, or `none`. Figure 2-2 shows some of these styles for the axis lines.




Figure 2-2: Left, center, and right axis lines.

Clearly, we need to use `axis lines = left` as a parameter to make the L-shaped frame most economic graphs have. In addition, the arrows at the end of the axes can be removed by appending an asterisk to the end of the `axis lines` so we use, `axis lines* = left`.

2.3 Axis ticks, scale, and clipping

Economic graphs typically do not have numbers running along the axes. This is because most of them are theoretical diagrams and apply irrespective of the orders of magnitude involved. To this end, we need to remove the lines and numbers running along the x and y-axes—which are called axis ticks and axis tick labels respectively. The parameters that do this are `xtick = {tick numbers}` and `ytick = {tick numbers}` for the x and y-axes respectively. Within the curly braces after the `xtick` and `ytick` parameters, we put a list of numbers separated by a comma which we want to run along the respective axes. Since economic graphs, for the most part, keep the zero at the bottom-left, we will keep it in the `xticks` as such `xtick = {0}`. Since we want no axis ticks on the y-axis, we write `ytick = \empty`. The case of an empty axis ticks numbers list is a special case. Because `pgfplots` would use the default axis ticks if we had nothing within the curly braces, we need to specify that it is empty with `\empty`. To make non-blank axis lines—for instance, graphs with specific, empirical values—Section 4: Empirical curves covers grid lines and axis ticks.

Considering the size of the figure on the page, the default size is slightly too small to present much illustration without cluttering the space available. To slightly increase the size of the graph, we write the parameter `\scale = 1.2`. Of course, if you have a more complex or important illustration, a greater factor can be used to create bigger graphs. On the other side, a lesser factor can be used to create smaller graphs. A scale factor of 1.2 occupies about a third of the page.

The last parameter we need to add is `clip = false`. Normally, `pgfplots` cuts off any plots outside the graph, which is useful for automatically bounding functions. However, we want to be able to put labels outside the graph, so we disable that feature with this parameter. Instead, we will be manually indicating the domain and range of plotted functions.

2.4 Axis labels

Finally, the last consideration we have to address before plotting anything on the graph are the axis labels. It is more convenient to put them at the ends of axes instead of conventionally at the bottom (for the x-axis) and to the left (for the y-axis) because we want to use that space for other things (such as variable labels and dimension lines). To put labels at the end of the x and y-axes, the commands take the form of:

```
% Labels
\node [right] at (current axis.right of origin) {$x$};
\node [above] at (current axis.above origin) {$y$};
```

As a reminder, semicolons should be at the end of these lines because they are commands as opposed to the parameters that have been used thus far.² For the demonstration, we will use the x and y-axis labels of x and y respectively. Putting the dollar signs around the x and y puts them in LaTeX's math mode (corresponding to x and y), allowing us to write the mathematical expressions LaTeX is known for.

After all of these considerations are addressed, we end up with Figure 2-3:

²Subsection 3.5: Labelling explains why this command works in more detail.




Figure 2-3: The finished x and y-axes, blank and L-shaped.

Now the graph is ready to have plots (e.g. functions) drawn upon it. The code that produces Figure 2-3 is as follows:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Labels
\node [right] at (current axis.right of origin) {$x$};
\node [above] at (current axis.above origin) {$y$};
\end{axis}
\end{tikzpicture}
\end{center}
```

3 Plotting, colours, and labelling (ex. indifference map and budget constraint)

With our axes, we can begin drawing on the graph. Suppose we wanted to illustrate why demand is downward-sloping—meaning that quantity demanded falls as price rises—using indifference curves. More specifically, we want to plot an indifference map between good A and B with budget constraints showing an increase in the price of good A corresponding to a decrease in the quantity demanded of the good.

3.1 Plotting functions

First, we want to draw an indifference curve defined by, say, $U_1 : B = \frac{10}{A^2} + 1$. This is done by the following command:

```
\addplot[domain = 0:10, restrict y to domain = 0:10, samples =  
400, color = red]{10/(x^2)+1};
```

This is a long command. Let us take this in pieces. The `addplot` parameter of `domain = min:max` restricts the function graphed to the domain bounded by the minimum x-value, *min* (which is 0 in this case), and maximum x-value, *max* (which is 10 in this case). Similarly, `restrict y to domain = min:max` restricts the function to the range bounded between *min* and *max* (which are 0 and 10 in this case). (Restricting y to a “domain” is a slight misnomer because the proper mathematical term for the set of y-values is “range”.) Note, the reason why we need to define the domain and range, as mentioned before, is that we disabled pgfplots’s clipping feature. The `addplot` parameter `samples = n` tells pgfplots to take a certain number of x-values, *n*, at which to evaluate the function and draw a line through all the calculated coordinates. In practice, the higher the number of samples, the smoother the plotted curve. The last `addplot` parameter, `color = colour name`, is self-explanatory. It colours the line to be the colour name. Colours are addressed in more detail in the next subsection. In Subsection 3.2: Colours, there is a list of the predefined colour names along with an explanation on how to add more predefined colour sets and define new colour names.

Moving on to the mathematical expression within the curly braces, this is where

we put the function's expression, $\frac{10}{x^2} + 1$. It is important to note that the conventional notation of juxtaposition (placing numbers and letters side by side) for multiplication throws an error. Instead, an asterisk—“*”—can only indicate multiplication. Otherwise, conventional computer notation applies, with addition, subtraction, division, and grouping represented by “+”, “-”, “/”, and “()” respectively. Plotting the aforementioned command yields Figure 3-1.




Figure 3-1: Plotting an indifference curve between good A and good B.

The code that produces Figure 3-1 is:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Indifference curve
```

```

\addplot[domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = red] {10/(x^2)+1};

% Labels
\node [right] at (current axis.right of origin) {$A$};
\node [above] at (current axis.above origin) {$B$};
\end{axis}
\end{tikzpicture}
\end{center}

```

One thing to notice about Figure 3-1 and its code is that the x-axis label is A and the y-axis label is B . This is because, in this example, we want to show that an increase in the price of good A causes a decrease in the quantity demanded of the good.

Naturally, an indifference map contains more than one curve. So, let us plot multiple curves by translating the first curve to the right by one unit and up by one unit. In transformation notation, this corresponds to $(x, y) \rightarrow (x + 1, y + 1)$. Transforming the first curve, $U_1 : B = \frac{10}{A^2} + 1$, results in the second curve, $U_2 : B = \frac{10}{(A-1)^2} + 2$. Repeating this transformation results in, $U_3 : B = \frac{10}{(A-2)^2} + 3$, $U_4 : B = \frac{10}{(A-3)^2} + 4$, and $U_5 : B = \frac{10}{(A-4)^2} + 5$. Plotting all five of these curves yields Figure 3-2.




Figure 3-2: Plotting an indifference map between good A and good B.

The indifference curves in Figure 3-2 are plotted by the following code snippet:

```
% Indifference curves
\addplot[domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/(x^2)+1};
\addplot[domain = 1:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-1)^2)+2};
\addplot[domain = 2:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-2)^2)+3};
\addplot[domain = 3:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-3)^2)+4};
\addplot[domain = 4:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-4)^2)+5};
```

Note that the domain restrictions are different for each function. U_1 is restricted to $A \in (0, 10)$, while U_2 is restricted to $A \in (1, 10)$, and U_3 is restricted to $A \in (2, 10)$, and so on. The reason for this is because each indifference curve's equation has different vertical asymptotes, with it being $A = 0$ for U_1 , $A = 1$ for U_2 , and $A = 2$ for U_3 , etc. Because we only care about the graph of the equation at the right-hand side of the vertical asymptote, we restrict the utility functions to the domain right of the vertical asymptote. Figure 3-3 illustrates this domain restriction.




Figure 3-3: Restricting the domain of an indifference curve.

Now that we have the indifference map plotted, we need to plot the changing budget constraints in response to a rise in the price of good A. In effect, this means

two diagonal lines: one from a positive B-intercept to a positive A-intercept, and another line from that B-intercept to a reduced positive A-intercept. This represents the change in the maximum number of good A that can be bought because of a price increase, while the maximum number of good B remains the same. In addition, because we want to show the change in the consumer decisions after a constriction in budget constraint, the two budget constraints must be tangent to indifference curves. We will use U_3 to be tangent to the original budget constraint at $A = 4.7$ and U_2 to be tangent to the new one at $A = 3.3$. The original and new budget constraints are calculated to be the equations $B = 9.16 - 1.02A$ and $B = 9.16 - 1.59A$ respectively (these are approximations of the actual tangents). Logistically, the equations of these budget constraints can be calculated or approximated beforehand, manually or with software such as WolframAlpha or Desmos.³ Using the same `addplot` command which we used previously to plot the lines of the indifference curves, we can plot the two budget constraints.

Lastly, to make the budget constraint lines stand out in the midst of the indifference map, we need to thicken them. This is done by adding the parameter `thick` to the `addplot`. The preset line thicknesses are `ultra thin`, `very thin`, `thin`, `thick`, `very thick`, or `ultra thick`—or we could define our own thickness with the `line width = width` parameter where `width` is in units of length. In total, we have the following code snippet for the budget constraint lines:

```
% Budget constraints
\addplot[domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = blue, thick]{9.16-1.02*x};
\addplot[domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = blue, thick]{9.16-1.59*x};
```

Plotting them on the graph yields Figure 3-4.

³For example, with WolframAlpha, the original budget constraint equation can be found by inputting "tangent at $x = 4.7$ for $y = 10/((x-2)^2+3)$ ", which outputs " $y = 9.14744 - 1.01611x$ ". Manually, this is done by letting $U_3(A) = \frac{10}{(A-3)^2} + 4$, calculating the slope at the point $U'_3(4.7) \approx -1.02$. Using the point-slope form—the form $y - y_1 = m(x - x_1)$ —the budget constraint equation is $B = U'_3(4.7)(A - 4.7) + U_3(4.7)$ which is approximately equal to $B = 9.16 - 1.02A$.




Figure 3-4: Plotting budget constraints on the graph.

3.2 Colours

In the previous subsection, we used two colours—red and blue—to draw functions. You might be wondering how many colours we can use, or how to define new colours. Remember, we called the `xcolor` package because it expands LaTeX’s limited base capacities to manipulate and use colours.




There are 19 default colours predefined by LaTeX:



red (■), green (■), blue (■), cyan (■), magenta (■), yellow (■), black (■),
gray (■), white (■), darkgray (■), lightgray (■), brown (■), lime (■), olive (■),
orange (■), pink (■), purple (■), teal (■), and violet (■).⁴

We already saw red and blue in the previous subsection. Although limited, the default colour palette is diverse enough for this guide.

⁴By the way, in case you are wondering how I am producing these coloured squares, I explain how in Subsection 6.2: Adding a legend. The command `\fcolorbox{black}{red}{\textcolor{red}{\rule{\fontcharht\font{X}}{\fontcharht\font{X}}}}` produces ■.

However, if these colours are too limited, xcolor comes with a variety of options that defines more colours. For example, the dvipsnames option defines 68 more colours:

Apricot (), Aquamarine (), Bittersweet (), Black (), Blue (), Blue-Green (), BlueViolet (), BrickRed (), Brown (), BurntOrange (), Cadet-Blue (), CarnationPink (), Cerulean (), CornflowerBlue (), Cyan (), Dandelion (), DarkOrchid (), Emerald (), ForestGreen (), Fuchsia (), Goldenrod (), Gray (), Green (), GreenYellow (), JungleGreen (), Lavender (), LimeGreen (), Magenta (), Mahogany (), Maroon (), Melon (), MidnightBlue (), Mulberry (), NavyBlue (), OliveGreen (), Orange (), OrangeRed (), Orchid (), Peach (), Periwinkle (), PineGreen (), Plum (), ProcessBlue (), Purple (), RawSienna (), Red (), RedOrange (), RedViolet (), Rhodamine (), RoyalBlue (), RoyalPurple (), RubineRed (), Salmon (), Sea-Green (), Sepia (), SkyBlue (), SpringGreen (), Tan (), TealBlue (), Thistle (), Turquoise (), Violet (), VioletRed (), White (), WildStrawberry (), Yellow (), YellowGreen (), and YellowOrange ().

It is important to note that these colour names are case-sensitive. Meaning that `color = violet` () is different from `color = Violet` (), and `color = VIOLET` will throw an error.

To add the dvipsnames colour palette, we call the xcolor package with the dvipsnames option, so, `\usepackage[dvipsnames]{xcolor}`. If the `usepackage` command is called after `pgfplots`'s, an error will be thrown. This is because TikZ calls xcolor by default and clashes with this new `usepackage` command. So make sure that the xcolor package is called before the TikZ package. Alternatively, if this does not work, then we could add the parameter `xcolor = {dvipsnames}` to the document class command. For us, we would use, `\documentclass[12pt, xcolor = {dvipsnames}]{article}`.

There are other palettes aside from dvipsnames, such as svgnames, which defines 151 colours, or x11names which defines 317 colours. In addition, new colours can be manually defined using the command `\definecolor{name}{model}{variable 1, variable 2, variable 3}`, where the model is either `rgb`, `HTML`, `cmyk`, among others. For example, a light tan colour can be defined by

`\definecolor {name}{rgb}{0.95, 0.95, 0.92} (□).`⁵

A lot can be done with colours in LaTeX. We can colour text, have background colours, and make the page colour different from the standard white (though that would be ill advised, considering the strain on printers). Most of these are outside the scope of this guide and so will not be covered. Resources are available at Section 8: Resources, which includes documentation and guides on the xcolor package.

3.3 Plotting dashed line segments

Returning to the indifference map and budget constraints graph we were working on, the next step is to draw dashed lines on it. Pgfplots has a command for line graphs, wherein it will connect all inputted coordinates sequentially with a line going through them. The coordinates of the original decision point—where the original budget constraint is tangent to an indifference curve—are calculated to be (4.7, 4.37). So the dashed line will go from (4.7, 0), the A-intercept, to (4.7, 4.37), the decision point, to (0, 4.37), the B-intercept. The command that does this is as follows:

```
\addplot[color = black, dashed, thick] coordinates {(4.7, 0) (4.7, 4.37) (0, 4.37)};
```

As seen, the coordinates the line goes through should be listed within the curly braces. Also, because we want the line to be dashed, we include the parameter `dashed`. The other two parameters, `color = black` and `thick` have already been covered in the previous subsections. They respectively colour the resulting line black and make its line width thicker. Repeating this process for the new decision point at (3.3, 3.9) yields Figure 3-5.

⁵Incidentally, this is what I use as the background colour of the code snippets in this guide.




Figure 3-5: Plotting dashed lines on the graph.

3.4 Plotting coordinate points

Conveniently, imposing black coordinate points on the graph is very similar to drawing a line. In fact, to make a scatter plot, pgfplots uses the same command that was previously used to draw connected line segments. Then two parameters are added. The first is `mark = *`, which draws a filled circle at every coordinate the line connects. Marks other than “*” can be used, such as “x”, “+”, “|”, “o”, (all of which look like the character representing it), and many more marks. Second, to remove the line, we write the parameter `only marks`. In addition, to increase the size of the marks, we add the parameter `mark size = 3`. In total, we have the following command to plot the coordinate points:

```
\addplot[color = black, mark = *, only marks, mark size = 3pt]
coordinates {(3.3, 3.9) (4.7, 4.37)};
```

Using this command on the graph yields Figure 3-6.




Figure 3-6: Plotting coordinate points on the graph.

3.5 Labelling

Finally, to finish the graph we need to label all the relevant parts of it. Specifically, we need to put letters near the A and B-intercepts, the coordinate points, and the ends of the functions. To label a part of the graph, the command takes the form of:

```
\node [position] at (coordinate) {label};
```

This command will place some text, the label, next to the coordinate specified. In addition, where the label ends up near the coordinate depends on the position, it can either be `left`, `right`, `above`, or `below` the coordinate. For instance, to put a label of Q_A below $(4.7, 0)$, where the dashed line extending downward from the original decision point meets the A-axis, we write the following command:

```
\node [below] at (4.7, 0) {$Q_A$};
```

We have already used a form of the node command previously when labelling our axes. Notice, however, that the coordinate point for an axis label is different because it refers to the end of an axis. Using the labelling command to label the rest of the graph, such as the functions and coordinates, we end up with Figure 3-7.




Figure 3-7: The effect of an increase in the price of good A on the budget constraint and indifference map between good A and good B.

Finally, we have completed the indifference map illustrating why demand slopes downward. It is because Q'_A is less than Q_A when the budget shifts from M to M' . More importantly, we illustrated this in LaTeX! Figure 3-7 is produced by the following code, combining all of the elements covered in this section:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Indifference curves
\addplot[domain = 0:10, restrict y to domain = 0:10, samples =
400, color = red]{10/(x^2)+1};
```

```

\addplot[domain = 1:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-1)^2)+2};
\addplot[domain = 2:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-2)^2)+3};
\addplot[domain = 3:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-3)^2)+4};
\addplot[domain = 4:10, restrict y to domain = 0:10, samples =
    400, color = red]{10/((x-4)^2)+5};

% Budget constraints
\addplot[domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = blue, thick]{9.16-1.02*x};
\addplot[domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = blue, thick]{9.16-1.59*x};

% Dashed lines
\addplot[color = black, dashed, thick] coordinates {(4.7, 0) (4.7,
    4.37) (0, 4.37)};
\addplot[color = black, dashed, thick] coordinates {(3.3, 0) (3.3,
    3.9) (0, 3.9)};

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
    coordinates {(3.3, 3.9) (4.7, 4.37)};

% Labels
\node [right] at (current axis.right of origin) {$A$};
\node [above] at (current axis.above origin) {$B$};
\node [above] at (9.1, 0) {$M$};
\node [above] at (6, 0) {$M^\prime$};
\node [below] at (4.7, 0) {$Q_A$};
\node [below] at (3.3, 0) {$Q_A^\prime$};
\node [left] at (0, 4.5) {$Q_B$};
\node [left] at (0, 3.7) {$Q_B^\prime$};
\node [above] at (5.1, 4.2) {$D$};
\node [above] at (2.9, 2.8) {$D^\prime$};
\node [right] at (10, 1.1) {$U_1$};
\node [right] at (10, 2.12) {$U_2$};

```

```

\node [right] at (10, 3.16) {$U_3$};
\node [right] at (10, 4.2) {$U_4$};
\node [right] at (10, 5.4) {$U_5$};
\end{axis}
\end{tikzpicture}
\end{center}

```

4 Empirical curves (ex. market equilibrium)

At the start of this guide, I assumed that we were working on a blank graph. Meaning, the x and y-axis did not have any ticks because they do not have any values associated with them. These types of graphs are useful for theoretical models—applicable to relevant situations no matter the order of magnitude. Sometimes, however, we actually have empirical data. In this case, we might want grid lines on which to impose plots such as functions, coordinates, dashed lines, and so on. A couple modifications are needed to convert our "blank" graph to service empirical values. Specifically, we want to impose grid lines and then work out the legibility problems such as distancing axis labels and putting a white background colour behind labels in the graph to avoid clutter. Let us transform a market equilibrium curve, Figure 4-1, so that is able to reflect empirical data.




Figure 4-1: The market equilibrium for the widget market (on a blank graph).

The elements in Figure 4-1 have been explained in previous sections. It is produced by the following code:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Supply and demand curves
\addplot[color = blue, very thick] coordinates {(1, 9) (9, 1)};
\addplot[color = red, very thick] coordinates {(1, 1) (9, 9)};

% Dashed lines
\addplot[color = black, dashed, thick] coordinates {(0, 5) (5, 5)
(5, 0)};
```

```

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
    coordinates {(5,5)};

% Labels
\node [right] at (current axis.right of origin) {Widgets,  $Q$ };
\node [above] at (current axis.above origin) {Price,  $P$ };
\node [above] at (5, 5.2) { $E$ };
\node [left] at (0, 5) { $P_E$ };
\node [below] at (5, 0) { $Q_E$ };
\node [right] at (9, 1) { $D$ };
\node [right] at (9, 9) { $S$ };
\end{axis}
\end{tikzpicture}
\end{center}

```

4.1 Major and minor grid lines

To start with, we want axis ticks and numbers to run along both axes. This was covered briefly in Subsection 2.3: Axis ticks, scale, and clipping, but there we were focused on removing them. As a reminder, the axis parameters which control the ticks are `xtick = {tick numbers}` and `ytick = {tick numbers}`. Instead of them containing only 0 or being empty, we want them to be the numbers which fit our data. In this case, the set of tick numbers will be $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$ for both of them. Next, we want to impose grid lines upon the graph. We are still working with axis parameters, not commands. Usefully, there are two types of grid lines, major grid lines and minor grid lines. Major grid lines extend from the axis ticks and numbers. Meanwhile minor grid lines extend from smaller ticks between major ticks and do not have a number associated with them. They both have different styles and so can aesthetically convey more precision. To show both grids, include the axis parameter `grid = both`. Similarly, if we want only one or none of the grids, we would use the value of `major`, `minor`, or `none`—though by default neither grid is shown. To set the number of minor ticks between major ticks, the axis parameter `minor tick num = number` is used,

where the number is the number of minor ticks. In this case, we want one minor tick halfway between the major ticks—otherwise it would be too cluttered—so the number of minor ticks is 1. Moreover, we want the major grid lines to be solid lines and the minor grid lines to be dotted lines. This way, half and full increments are visually evident. To change the style of major and minor grid lines to solid and dotted lines, we use `grid style = solid` and `minor grid style = dotted` respectively. More styles include `dashed`, `thin`, and `thick`. Defining the grid style to be solid is redundant since that is the default, but I have opted to make it explicit. So far we have added, or modified, the parameters in the following code snippet:

```
xtick = {0,1,2,3,4,5,6,7,8,9,10},  
ytick = {0,1,2,3,4,5,6,7,8,9,10},  
grid = both,  
minor tick num = 1,  
grid style = solid,  
minor grid style = dotted,
```

This snippet adds axis ticks from 1 to 10 in increments of 1 on both the x and y axes. In addition, it shows both the minor and major grid lines, sets the number of minor ticks between major ticks to one tick, and makes major grid lines solid lines and minor grid lines dotted lines. Using these parameters to the blank market equilibrium graph results in Figure 4-2.




Figure 4-2: Imposing grid lines on the graph.

Ah! That is ugly! While we got the axis ticks and grid lines we were seeking, this caused other issues for the legibility of this graph. Specifically, the labels of Q_E and P_E need to be moved away from the axis numbers, the axis labels should be moved slightly away from the graph, and the labels of E , S , and D need to have white backgrounds to reduce clutter. In addition, there are two zeros at the origin—one is enough.

4.2 Distancing labels, colouring label backgrounds

Starting with the former issue, the x and y-intercept labels need to be moved away from the graph to make room for the axis numbers. Recall that labels take the form of:

```
\node [position] at (coordinate) {label};
```

While the position does a good job in most situations of determining how far the label should be from the coordinate, sometimes we want to increase this distance. In such situations, the distance can be specified with `position = distance`. So,

for example, the position for Q_E can be changed to `below = 10pt`. Repeating this for the P_E label, we have the following code snippet:

```
\node [left = 10pt] at (0, 5) {$P_E$};
\node [below = 10pt] at (5, 0) {$Q_E$};
```

Moving on to the latter issues. Getting rid of the y-axis label of 0 and moving the axis labels can both be done by changing the domain and range of the graph shown. To get rid of the 0 on the y-axis, we define the minimum of the range to be 0.01 (any very small number suffices, though this varies depending on the order of magnitude of the data). So we have `ymin = 0.01` as an axis parameter. To distance the axis labels, the maximums of both the domain and range should be extended to 10.5 (this too varies depending on the order of magnitude). In sum, we have the following code snippet:

```
xmin = 0, xmax = 10.5,
ymin = 0.01, ymax = 10.5,
```

Finally, we want to colour the background of labels on the graph white. Otherwise, we would be stuck with lines cutting through E , S , and D , which does not look good. This also involves a relatively simple adjustment. We add the parameter `fill = white` to the node commands of the labels. The code snippet for the changed labels are:

```
\node [above = 5pt, fill = white] at (5, 5.2) {$E$};
\node [right, fill = white] at (9, 1) {$D$};
\node [right, fill = white] at (9, 9) {$S$};
```

The label of E has also been distanced by 5 pt using the positioning parameter as explained in the preceding paragraph. This was done to avoid covering the supply and demand lines with the white background. There are other ways to resolve this inconvenience. For example, because pgfplots layers these plots sequentially—where later commands cover previous ones—simply having the node command for the label above the ones for the supply and demand line curves could work. That way, the commands for the supply and demand curves would be executed after the label and would cover up the white background.

Combining all of these adjustments, we end up with Figure 4-3.




Figure 4-3: The market equilibrium for the widget market.

And so we have finished adjusting the blank graph to one with ticks and grid lines so that it can accommodate empirical data and represent specific values. The code for Figure 4-3 is:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10.5,
ymin = 0.01, ymax = 10.5,
axis lines* = left,
xtick = {0,1,2,3,4,5,6,7,8,9,10},
ytick = {0,1,2,3,4,5,6,7,8,9,10},
grid = both,
minor tick num = 1,
grid style = solid,
minor grid style = dotted,
clip = false,
]
```

```

% Supply and demand curves
\addplot[color = blue, very thick] coordinates {(1, 9) (9, 1)};
\addplot[color = red, very thick] coordinates {(1, 1) (9, 9)};

% Dashed lines
\addplot[color = black, dashed, thick] coordinates {(0, 5) (5, 5)
(5, 0)};

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
coordinates {(5,5)};

% Labels
\node [right] at (current axis.right of origin) {Widgets,  $Q$ };
\node [above] at (current axis.above origin) {Price,  $P$ };
\node [above = 5pt, fill = white] at (5, 5.2) { $E$ };
\node [left = 10pt] at (0, 5) { $P_E$ };
\node [below = 10pt] at (5, 0) { $Q_E$ };
\node [right, fill = white] at (9, 1) { $D$ };
\node [right, fill = white] at (9, 9) { $S$ };
\end{axis}
\end{tikzpicture}
\end{center}

```

5 Shading, arrows, and graphs side by side (ex. supply increase, price-taking firm)

A lot can be done by plotting curves, lines, and coordinate points. But sometimes we need to point out features or indicate areas on the graph. This can be done by drawing arrows and shading in areas on the graph with colour. Let us use these tools to illustrate the bumper harvest paradox: why farmers' incomes fall after a bountiful harvest. A bumper harvest corresponds to a supply shift rightward in the market equilibrium graph where demand is relatively price-inelastic. This is depicted in Figure 5-1.




Figure 5-1: The effect of a bumper harvest on the farm product market.

Figure 5-1 is graphed using the tools already explained in the sections above. The code is as follows:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Supply and demand curves
\addplot[color = blue, very thick] coordinates {(5,9) (8,1)};
\addplot[color = red, very thick] coordinates {(3,1) (6,9)};
\addplot[color = red, very thick] coordinates {(6,1) (9,9)};

% Dashed lines
```

```

\addplot[color = black, dashed, thick] coordinates {(0, 7.67) (5.5,
    7.67) (5.5, 0)};
\addplot[color = black, dashed, thick] coordinates {(0, 3.67) (7,
    3.67) (7, 0)};

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
    coordinates {(5.5, 7.67) (7, 3.67)};

% Labels
\node [right] at (current axis.right of origin) {Farm products,  $Q$ 
    $};
\node [above] at (current axis.above origin) {Price,  $P$ };
\node [right] at (5.5, 7.67) { $E$ };
\node [right] at (7, 3.67) { $E^{\prime}$ };
\node [right] at (8, 1) { $D$ };
\node [above] at (6, 9) { $S$ };
\node [right] at (9, 9) { $S^{\prime}$ };
\node [left] at (0, 7.67) { $P_E$ };
\node [left] at (0, 3.67) { $P^{\prime}$ };
\node [below] at (5.5, 0) { $Q_E$ };
\node [below] at (7, 0) { $Q^{\prime}$ };
\end{axis}
\end{tikzpicture}
\end{center}

```

5.1 Opacity and transparency

Right now, both supply curves are the same opaque red. To make it evident which is the original supply curve and which is the new one, we will make the new one slightly transparent. We do this by adding the parameter `opacity = opacity` value to the `addplot` command, where the opacity value is the percent the plot is opaque ranging from 0, perfectly transparent, to 1, perfectly opaque. Making the new curve mostly transparent seems to make it evident that it is new. So we write `opacity = 0.3` for its `addplot` command. The result of making the new supply curve more transparent is shown in Figure 5-2.




Figure 5-2: Making the new supply curve on the graph more transparent.

It is important to note that the `xcolor` package—which was mentioned in Sub-section 3.2: Colours—offers a more concise though slightly less readable way of making colours transparent. When using a colour, an exclamation mark can be used after the colour name, right after which comes the opacity value—“colour name!opacity value”. This one ranges from 0 being perfectly transparent to 100 being perfectly opaque. For example, `red!10` (□) is mostly transparent, `red!50` (□) is half transparent, and `red!90` (□) is mostly opaque. (Although, if I might add, that “mostly transparent red” looks more like salmon.)

5.2 Shading in boxes with colours

Now that we have distinguished which supply curve is the new one (more than just using the label S'), we also want to indicate the areas that represent the total revenue loss and the total revenue gain for the farmers. That way, we can graphically show that the loss in total revenue is much greater than the total revenue gain. The command that fills in an area on the graph with colour takes the form of:


```
\fill[colour name] (A) -- (B) -- (C) -- ...;
```

In this command, (A), (B), (C), ..., are an indefinite number of coordinates. The command colours in the polygon defined by the coordinates with the specified colour. Looking at the graph, the original equilibrium point is (5.5, 7.67) and the new equilibrium point is (7, 3.67). Meaning that, for the loss area, we need to colour in orange the square defined by the coordinates (0, 7.67), (0, 3.67), (5.5, 3.67), and (5.5, 7.67). Similarly, we shade green the gain area defined by the coordinates (5.5, 0), (5.5, 3.67), (7, 3.67), and (7, 0). So far, we have the following code snippet:

```
\fill[orange] (0, 7.67) -- (0, 3.67) -- (5.5, 3.67) -- (5.5, 7.67)
;
\fill[green] (5.5, 0) -- (5.5, 3.67) -- (7, 3.67) -- (7, 0);
```

Since we want these coloured areas to be unobtrusive, the colours need to be mostly transparent. This can be done the same way it is done to plots as shown in the previous subsection, Subsection 5.1: Opacity and transparency, either through TikZ's opacity parameter or through xcolor's exclamation mark notation. So the snippet the graph will be using is:

```
\fill[orange, opacity = 0.1] (0, 7.67) -- (0, 3.67) -- (5.5, 3.67)
-- (5.5, 7.67);
\fill[green, opacity = 0.1] (5.5, 0) -- (5.5, 3.67) -- (7, 3.67)
-- (7, 0);
```

Unfortunately, we cannot simply put these commands into the axis environment. There are a couple considerations to address so that these coloured areas do not interfere and overlap with other features on the graph. Because pgfplots processes commands sequentially so that later plots are layered upon earlier plots, the fill commands should precede every other one—otherwise the coloured areas could cover up another part of the graph. Also, since fill commands cannot precede the axis parameters, the axis parameter `axis on top` should be added. This ensures that the axis lines are not covered up by plots and other commands. The coloured areas are shown in Figure 5-3.




Figure 5-3: Colouring in areas on the graph.

5.3 Straight and curly arrows

While this graph illustrates the bumper harvest through the relative sizes of the shaded areas, it is missing information about what those shaded areas represent. Of course, a reader could infer from context that the orange area represents the total revenue loss and the green area represents the total revenue gain, but we can also make the graph more readable for them. So let us label the shaded areas with curly arrows. Also, since we want to illustrate the supply curve moving from the original to the new one, we can draw a straight arrow from the original to the new curves.

Starting with the latter, since that is easier, the command for drawing arrows takes the form of:

```
\draw[-arrowhead] (A) to (B);
```

This command draws an arrow starting from (A) and ending at (B). The tip of the arrow at B is specified by the `arrowhead` value. TikZ has a limited array of arrowheads to choose from, so that is why we loaded the `arrows.meta` TikZ library in Subsection 1.3: Packages. With this library, we have access to more

arrowheads such as `Triangle`, `Circle`, `Square`, `Diamond`, and so on, all of which are self-descriptive. The `Triangle` arrowhead is the most appropriate to show the shift from the original to the new supply curves.

Using this command, we will draw an arrow from the right of the original supply curve to the left of the new supply curve. This corresponds to an arrow from (6.3, 8.5) to (8.3, 8.5) with an arrowhead of `Triangle`. So we have the following code snippet:

```
\draw[-Triangle] (6.3, 8.5) to (8.3, 8.5);
```

While this will do the job, there are a couple changes we should make to improve the aesthetics and readability of the arrow. First, it would be better if the arrow was a mostly transparent red to show it is modifying the supply curve. Conveniently, this is just like filling areas with colours and colouring plots as we have done in before. We write `red, opacity = 0.3` as parameters to the draw command for the arrow. Second, arrowheads are small by default. We would like to make the supply curve shift arrow slightly bigger. This is slightly harder to do. To begin with, we surround the arrowhead value with curly braces to indicate that we are using a custom arrowhead. So we have `-{Triangle}`. Then, as parameters *to the arrowhead value* we add the parameters `length = 4mm, width = 2mm` which modify the dimensions of the arrowhead. Accordingly, we have the final code snippet for the supply shift arrow:

```
\draw[-{Triangle[length = 4mm, width = 2mm]}, red, opacity = 0.3]
(6.3, 8.5) to (8.3, 8.5);
```

Using this command to draw the arrow upon the graph results in Figure 5-4.




Figure 5-4: Drawing a straight line from the original supply curve to the new one.

Next, we will label the gain and loss areas. Simply putting a vertical arrow from below the x-axis to the green area would be too cluttered. We would have the labels Q , Q' , and an arrow all sharing that small space. Instead, let us draw a curly arrow pointing from the bottom-right of the axis to the centre of the green area curving in a counterclockwise direction so that it avoids the Q and Q' labels. Curly arrows are drawn with a draw command in the form of:

```
\draw[-arrowhead] (A) to [out = out angle, in = in angle] (B);
```

Notice that this is nearly identical to the command for straight arrows, except for one addition. Within the square brackets after the “to”, the parameters specify the angles from the starting coordinate and to the ending coordinate which the line for the curve should be drawn. They are given in arc degrees, meaning from 0 for right, to 90 for up, to 180 for left, to 270 for down, to 360 for right again, and everything in between. I have found beforehand that the starting coordinate for the curly arrow is (10, -1.5) and the ending coordinate is (6.5, 0.7). Since the arrow should spring forth upward from the starting coordinate and comes into the ending coordinate from the right, the out angle is 90 and the in angle is 0. And so the code snippet

producing the curly arrow pointing at the gain area is:

```
\draw[-Triangle] (10, -1.5) to [out = 90, in = 0] (6.5, 0.7);
```

Repeating this for the loss area, for an arrow from the top to the orange box, and adding text labels (which was covered in Subsection 3.5: Labelling), we have this code snippet:

```
% Labels
\node [right] at (5, 11) {Total revenue lost};
\node [above] at (10, -3) {Total revenue gained};

% Arrows
\draw[-Triangle] (5, 11) to [out = 180, in = 90] (2, 6);
\draw[-Triangle] (10, -1.5) to [out = 90, in = 0] (6.5, 0.7);
```

Combining the straight and curly arrows, Figure 5-5 is produced.




Figure 5-5: The bumper harvest paradox: why farmers' revenues decrease after a large harvest.

Figure 5-5 is produced by the following code:

```

\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
axis on top,
clip = false,
]
% Colouring areas
\fill[orange, opacity = 0.1] (0, 7.67) -- (0, 3.67) -- (5.5, 3.67)
-- (5.5, 7.67);
\fill[green, opacity = 0.1] (5.5, 0) -- (5.5, 3.67) -- (7, 3.67)
-- (7, 0);

% Supply and demand curves
\addplot[color = blue, very thick] coordinates {(5,9) (8,1)};
\addplot[color = red, very thick] coordinates {(3,1) (6,9)};
\addplot[color = red, opacity = 0.3, very thick] coordinates
{(6,1) (9,9)};

% Dashed lines
\addplot[color = black, dashed, thick] coordinates {(0, 7.67) (5.5,
7.67) (5.5, 0)};
\addplot[color = black, dashed, thick] coordinates {(0, 3.67) (7,
3.67) (7, 0)};

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
coordinates {(5.5, 7.67) (7, 3.67)};

% Labels
\node [right] at (current axis.right of origin){Farm products, $Q
$};

```

```

\node [above] at (current axis.above origin) {Price,  $\$P$ };
\node [right] at (5.5, 7.67) { $\$E$ };
\node [right] at (7, 3.67) { $\$E^{\prime}$ };
\node [right] at (8, 1) { $\$D$ };
\node [above] at (6, 9) { $\$S$ };
\node [right] at (9, 9) { $\$S^{\prime}$ };
\node [left] at (0, 7.67) { $\$P_E$ };
\node [left] at (0, 3.67) { $\$P^{\prime}$ };
\node [below] at (5.5, 0) { $\$Q_E$ };
\node [below] at (7, 0) { $\$Q^{\prime}$ };
\node [right] at (5, 11) {Total revenue lost};
\node [above] at (10, -3) {Total revenue gained};

% Arrows
\draw[-{Triangle[length = 4mm, width = 2mm]}, red, opacity = 0.3]
    (6.3, 8.5) to (8.3, 8.5);
\draw[-Triangle] (5, 11) to [out = 180, in = 90] (2, 6);
\draw[-Triangle] (10, -1.5) to [out = 90, in = 0] (6.5, 0.7);
\end{axis}
\end{tikzpicture}
\end{center}

```

5.4 Arranging graphs side by side

In economics, it is common to arrange the market equilibrium graph to be side by side with the price-taking firm graph. This is useful to show how shifts in the market sets the price with which individual firms need to operate. More generally, we will use this example to demonstrate how we can place graphs beside each other. Suppose we wanted to show how the downward price movement in the market equilibrium causes a decrease in the total economic profit for individual farmers. The price-taking graph of this scenario is shown in Figure 5-6.




Figure 5-6: The effect of a bumper harvest on individual farmers.

The above graph is produced entirely with the tools and techniques already covered in previous sections. It is a combination of plots, coordinate points, labels, colours and transparency, and so on. Figure 5-6 is produced by the following code:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
axis on top,
clip = false,
]
```



```

% Colouring areas
\fill[orange, opacity = 0.1] (0, 3.67) -- (3.68, 3.67) -- (3.68,
    7.67) -- (0, 7.67);
\fill[green, opacity = 0.1] (3.68, 1) -- (5.45, 1) -- (5.45, 3.67)
    -- (3.68, 3.67);
\fill[green, opacity = 0.1] (0, 1) -- (0, 2) -- (3.68, 2) --
    (3.68, 1);

% Curves
\addplot [domain = 1:10, restrict y to domain = 0:10, samples =
    400, color = blue, very thick] {(x-1)^2+0.5};
\addplot [domain = 1:10, restrict y to domain = 0:10, samples =
    400, color = blue, opacity = 0.3, very thick] {(0.4*(x-1))
    ^2+0.5};
\addplot [domain = 1:10, restrict y to domain = 0:10, samples =
    400, color = red, very thick] {(0.55*(x-5))^2+1.5};
\addplot [domain = 1:10, restrict y to domain = 0:10, samples =
    400, color = red, opacity = 0.3, very thick] {(0.15*(x-5))^2+1
    };
\addplot [domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = teal, very thick] {7.67};
\addplot [domain = 0:10, restrict y to domain = 0:10, samples =
    400, color = teal, opacity = 0.3, very thick] {3.67};

% Dashed lines
\addplot[color = black, dashed, thick] coordinates {(3.68, 0)
    (3.68, 7.67)};
\addplot[color = black, dashed, thick] coordinates {(5.45, 0)
    (5.45, 3.67)};
\addplot[color = black, dashed, thick] coordinates {(0, 1) (5.45,
    1)};
\addplot[color = black, dashed, thick] coordinates {(0, 2) (3.68,
    2)};

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
    coordinates {(3.68, 7.67) (5.45, 3.67)};

```

```

% Labels
\node [right] at (current axis.right of origin) {Farm products,  $q$ 
    $};
\node [above] at (current axis.above origin) {Price,  $p$ };
\node [below right] at (3.68, 7.67) {$e$};
\node [above left] at (5.45, 3.67) {$e^{\prime}$};
\node [above] at (4, 10) {$MC$};
\node [above] at (10, 9) {$ATC$};
\node [right] at (10, 7.8) {$MR$};
\node [above] at (8.7, 10) {$MC^{\prime}$};
\node [right] at (10, 1.56) {$ATC^{\prime}$};
\node [right] at (10, 3.67) {$MR^{\prime}$};
\node [left] at (0, 7.67) {$P_E$};
\node [left] at (0, 3.67) {$P^{\prime}$};
\node [below] at (3.68, 0) {$q_e$};
\node [below] at (5.45, 0) {$q^{\prime}$};
\node [above] at (2, 12) {Total economic profit lost};
\node [above] at (3, -3) {Total economic profit gained};

% Arrows
\draw[-Triangle] (3, -1.5) to [in = 180, out = 90] (5, 2);
\draw[-Triangle] (3, -1.5) to (3, 1.7);
\draw[-Triangle] (2, 12) to (2, 6.5);
\end{axis}
\end{tikzpicture}
\end{center}

```

We want it arranged such that the market equilibrium graph (Figure 5-5) is on the left-hand side of a panel while the price-taking firm (Figure 5-6) is on the right-hand side. Surprisingly, this is actually very simple. All we need to do is create another axis environment right after the first one, but still within the same tikzpicture environment. This will overlay the two graphs on top of each other. Then, we add the parameter `shift = (axis cs: x-shift, y-shift)` to the second axis environment where the x and y-shifts are the amounts by which the second graph will be moved away from its original position. Since we want to shift the second graph (the price-taking firm) to the right of the first graph (the market equilibrium), we will use an x-shift of 17 and a y-shift of 0. So we have:

```

\begin{center}
\begin{tikzpicture}
\begin{axis}[
% Insert parameters for Graph 1
]
% Insert commands for Graph 1
\end{axis}
\begin{axis}[
% Insert parameters for Graph 2
shift = {(axis cs: 17, 0)},
]
% Insert commands for Graph 2
\end{axis}
\end{tikzpicture}
\end{center}

```

Finally, we need to make an adjustment to fit the figure on the page. Since the figure's width is greater than what the page margins allow, we tell TikZ to disregard the page margins by widening the picture frame. To do this, we write `\hspace*{-3cm}` before the `\begin{tikzpicture}` and after the `\end{tikzpicture}` which widens the picture's frame to the left by 3 cm and to the right by 3 cm respectively. In total, putting all of this together yields Figure 5-7.




Figure 5-7: The effect of a bumper harvest on the farm product market (L) and individual farmers (R).

The code that produces Figure 5-7 is:

```
\begin{center}
\hspace*{-3cm}\begin{tikzpicture}
\begin{axis}[
% Insert parameters for Graph 1
]
% Insert commands for Graph 1
\end{axis}
\begin{axis}[
% Insert parameters for Graph 2
shift = {(axis cs: 17, 0)},
]
% Insert commands for Graph 2
\end{axis}
\end{tikzpicture}\hspace*{-3cm}
\end{center}
```

The code is executed using the commands and parameters that produce Figures 5-5

and 5-6 for Graphs 1 and 2 respectively.⁶

5.5 Shading in the area under a curve

Closing off this section, we will briefly discuss how to shade the area under a curve—and by extension, how to shade in area between two curves—since that was not covered in the previous example. Suppose we wanted to shade in the area under a production-possibility frontier between good A and good B with a transparent blue in order to illustrate the output possibilities that exist which are not necessarily efficient. The production-possibility frontier is illustrated in Figure 5-8.




Figure 5-8: A production-possibility frontier between A and B (unshaded).

The code that produces Figure 5-8 is:

```
\begin{center}
\begin{tikzpicture}
```

⁶I will not include the full code for Figure 5-7 as it is too long and repeats the code snippets that were shown for Figures 5-5 and 5-6. Instead, the code that produces Figure 5-7 can be accessed at the GitHub repository: <https://github.com/jackypacky/pgf-econ-graphs>.

```

\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Production-possibility frontier
\addplot [domain = 0:10, restrict y to domain = 0:10, samples =
10000, color = blue, very thick] {(49-x^2)^0.5};

% Labels
\node [right] at (current axis.right of origin) {$A$};
\node [above] at (current axis.above origin) {$B$};
\end{axis}
\end{tikzpicture}
\end{center}

```

To shade in areas between two functions, we use a command that takes the form of:

```

\addplot [parameters] fill between [of = f and g];

```

where `parameters` are the parameters which modify the `addplot` command, and `f` and `g` are the *name paths* of the functions. For the parameters, we will use `blue`, and `opacity = 0.1` which creates a transparent blue similar to what we did in Subsection 5.1: Opacity and transparency.

To input the name paths `f` and `g`, however, we first need to define the name paths beforehand. To define the name path of a curve, we add the parameter `name path = label` to the `addplot` command where `label` is the name we want to assign it. So we will use the `addplot` parameter `name path = frontier` for the command which plots the production-possibility curve. Since we need another function between which we can shade in the area, and since we want to shade the area under the curve, we need to plot a new curve defined by the equation $y = 0$ and name it `axis`. For this new axis curve, we add the `addplot` parameter `line width = 0pt` to make it invisible—since we do not want it to show up on the

graph.

In total, shading in the area under the production-possibility curve is shown in Figure 5-9.




Figure 5-9: A production-possibility frontier between A and B.

Figure 5-9 is produced by the following code:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Production-possibility frontier
\addplot [domain = 0:10, restrict y to domain = 0:10, samples =
10000, color = blue, very thick, name path = frontier] {(49-x
^2)^0.5};
```

```

\addplot [domain = 0:10, restrict y to domain = 0:10, line width =
    0pt, name path = axis] {0};

% Colouring areas
\addplot [blue, draw = none, fill opacity = 0.1] fill between [of
    = frontier and axis];

% Labels
\node [right] at (current axis.right of origin) {$A$};
\node [above] at (current axis.above origin) {$B$};
\end{axis}
\end{tikzpicture}
\end{center}

```

6 Axis dimension lines, line breaks in labels, and a legend (ex. excise tax)

Continuing our discussion on enhancing the readability of graphs, let us now impose dimension lines and a legend upon a graph. Instead of putting a label on a graph pointing to the coloured areas, we could have instead added a legend showing what each area represents. In addition, to indicate what the differences between two axis labels are and mean, dimension lines can be used. Dimension lines are lines used in graphs, figures, and drawings to indicate sizes, quantities, and distances—dimensions. At each of the ends of dimension lines, there is a perpendicular line joined to it like a “T”.⁷

We will use an excise tax graph to illustrate the usefulness of dimension lines. Suppose we wanted to illustrate the tax incidence of an excise tax on gasoline. This is shown in Figure 6-1.

⁷Dimension lines kind of look like this: \perp — \perp




Figure 6-1: The effect of a gasoline excise tax.

The code that produces Figure 6-1 is:

```
\begin{center}
\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Colouring areas
\fill[teal, opacity = 0.1] (0, 3.53) -- (4.14, 3.53) -- (4.14,
5.96) -- (0, 5.96);
\fill[violet, opacity = 0.1] (0, 3.53) -- (4.14, 3.53) -- (4.14,
3) -- (0, 3);
\fill[teal, opacity = 0.25] (4.14, 3.53) -- (5.05, 3.53) -- (4.14,
5.96);
\fill[violet, opacity = 0.25] (4.14, 3) -- (5.05, 3.53) -- (4.14,
```

```

3.53);

% Supply and demand curves
\addplot[color = blue, very thick] coordinates {(3,9) (6,1)};
\addplot[color = red, very thick] coordinates {(1,1) (9,6)};
\addplot[color = red, opacity = 0.3, very thick] coordinates
    {(1,4) (9,9)};

% Dashed lines
\addplot[color = black, dashed, thick] coordinates {(0, 5.96)
    (4.14, 5.96) (4.14, 0)};
\addplot[color = black, dashed, thick] coordinates {(0, 3.53)
    (5.05, 3.53) (5.05, 0)};
\addplot[color = black, dashed, thick] coordinates {(0, 3) (4.1, 3)
    };

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
    coordinates {(4.14, 5.96) (5.05, 3.53)};

% Labels
\node [right] at (current axis.right of origin){Gasoline,  $\$Q$ };
\node [above] at (current axis.above origin) {Price,  $\$P$ };
\node [above] at (5.25, 3.6)  $\$E$ ;
\node [above] at (4.35, 6.1)  $\$E^{\prime}$ ;
\node [right] at (6, 1)  $\$D$ ;
\node [right] at (9, 6)  $\$S$ ;
\node [right] at (9, 9)  $\$S^{\prime}$ ;
\node [left] at (0, 3.7)  $\$P_E$ ;
\node [left] at (0, 2.8)  $\$P_S$ ;
\node [left] at (0, 5.96)  $\$P^{\prime}$ ;
\node [below] at (5.05, 0)  $\$Q_E$ ;
\node [below] at (4.14, 0)  $\$Q^{\prime}$ ;
\node [above] at (2, 4.7)  $\$A$ ;
\node [above] at (1, 2)  $\$B$ ;
\node [above] at (4.5, 3.5)  $\$C$ ;
\node [above] at (6.2, 2.2)  $\$F$ ;

```

```
% Arrows
\draw[-{Triangle[length=4mm, width=2mm]}, red, opacity = 0.3] (8,
    5.8) to (8, 7.8);
\draw[-Triangle] (1.3, 2.4) to [out = 0, in = 270] (2, 3.35);
\draw[-Triangle] (5.9, 2.6) to [out = 180, in = 270] (4.3, 3.4);
\end{axis}
\end{tikzpicture}
\end{center}
```

6.1 Adding dimension lines and line breaks in labels

Dimension lines are added in an identical manner to which arrows are drawn. Recall that the command for an arrow takes the form of:

```
\draw[-arrowhead] (A) to (B);
```

What I did not tell you was that arrowheads can be appended to both the end of an arrow and the start of an arrow. So the arrow with an arrowhead at its tail and head takes the form of:

```
\draw[out arrowhead-in arrowhead] (A) to (B);
```

Dimension lines use the arrowhead “|”, which is a nice symbolic description of how the arrow looks. We have:

```
\draw[|-|] (A) to (B);
```

The difference between Q_E and Q' is the reduction of the quantity of gas consumed after the imposition of the excise tax. To illustrate that on the graph, we will draw a dimension indicator from the bottom of Q' to Q_E . Below this, we will add a label with the text “Consumption reduction”. Since Q' is at $Q = 4.14$ and Q_E is at $Q = 5.05$, we have:

```
\draw[|-|] (4.14, -1) to (5.05, -1);
\node [below] at (4.59, -1) {Consumption reduction};
```

We would like to repeat this for the consumer incidence—between P_E and P' —and the producer incidence—between P_S and P_E . However, simply copying the above would not work because the label width would be too large. Instead, we need a way to insert a line break between the two words in both labels. This is

done by adding the parameter `align = left` to tell how the stacked text should align (either left, right, or center), and by adding two backslashes—“\\”—where we want the line to break. Accordingly, the code for dimension lines and labels for the consumer and producer incidences is:

```
\draw[|-|] (-1, 5.96) to (-1, 3.8);
\node [below, align = left] at (-2.3, 6) {Consumer \\ incidence};
\draw[|-|] (-1, 3.6) to (-1, 2.8);
\node [below, align = left] at (-2.3, 3.5) {Producer \\ incidence};
```

Adding a dimension line from P_S to P' representing the total tax and adding all of the dimension lines and labels mentioned so far results in Figure 6-2.




Figure 6-2: Adding dimension lines to the graph.

6.2 Adding a legend

Using a legend, we can specify what the areas A , B , C , F represent on the graph. Adding a legend combines a couple of the commands and parameters used in the

last subsection. While pgfplots has its own inbuilt legend feature, it is rather unwieldy to use. Instead, we are going to use a text label and put it at the top-right of the graph. The label should allow line breaks and have a black border drawn around it. So we will use:

```
\node [below right, draw, align = left] at (10.5, 10) {  
% Insert text here  
};
```

Recall that adding the parameter `align` to the `node` command allows line break—using `\\`—to be used within the label. Also, since the `align` is set to `left`, text lines will be flushed to the left. These were covered in the previous subsection. What is new, however, is the `draw` node parameter, which draws a black border around the label. In the command, the top-left of the legend will be at (10.5, 10), which corresponds to the top-right of the graph—where we want the legend to go. Adding text that describes what the areas represent, we have:

```
\node [below right, draw, align = left] at (10.5, 10) {  
A + C : Consumer surplus loss \\  
B + F : Producer surplus loss \\  
A + B : Total tax revenue \\  
C + F : Deadweight loss  
};
```

Using the above node command to add a legend to the graph, we get Figure 6-3.




Figure 6-3: Adding a legend to the graph.


As it stands, the legend looks a bit boring. To improve the aesthetics of the legend, we could add a colored background to the letters A, B, C, and F, to visually connect it to the coloured areas which they point to. It would look good if A were coloured a very transparent teal, B were coloured a very transparent violet, and so on.




Fortunately, the xcolor package has a command to do exactly that. The xcolor command to add a coloured background to text takes the form of:

```
\fcolorbox{frame colour}{background colour}{text}
```



In this command, the frame colour is the colour which surrounds the edges of the coloured background. The background colour is the colour behind the text. For the “A” text in the legend, since we want its background to be a very transparent teal, the command we want is:

```
\fcolorbox{black}{teal!10}{A}
```


The above command produces: . (As seen here, coloured text backgrounds can be used within inline text as well.) Notice that the command does not end with a semicolon, this is because it is not a TikZ command.

While the  looks fine, the widths of the background colour is determined by the widths of the text within it. So we hit a snag when using multiple background colours for texts with varying widths. For example, the background colour for  is wider than the one for . Beside each other, the inconsistent widths are noticeable. To fix this, we will use the `makebox` LaTeX command (`\makebox[width]{text}`) with the width of the height of “X”, (`\fontcharht\font`X`). Meaning the command we will be using to make a coloured background for the letters is:

```
\fcolorbox{black}{teal!10}{\makebox[\fontcharht\font`X]{A}}
```

which is quite overwhelming to look at without processing it in parts. You may be unfamiliar with the grave accent, “`”. On a conventional US keyboard, it is located at the top-left. It is used in LaTeX for an assortment of purposes, including quotation marks and adding a grave accent to characters.⁸ This way,  is just as wide as . Unfortunately, controlling heights is much harder and I will skip it in this guide. If we wanted to just have a coloured box with no text in it, we would use:

```
\fcolorbox{black}{teal!10}{\textcolor{teal!10}{\rule{\fontcharht\font`X}{\fontcharht\font`X}}}
```

Running the above command produces . This works because `\rule{width}{height}` generates a rectangular blob of ink. Incidentally, this is how I have been presenting the colours on their own throughout this guide. Finally, because adding a legend often makes a graph’s width bigger than what the page margins allow, we add `\hspace*{-3cm}` before `\begin{tikzpicture}` and after `\end{tikzpicture}`. We also did this in Subsection 5.4: Arranging graphs side by side when we encountered the same problem.

In total, we have the following code snippet for the legend:

⁸For example, adding a grave accent to “a” using `\`a` produces “à”.

```

% Legend
\node [below right, draw, align = left] at (10.5, 10) {
\colorbox{black}{teal!10}{\makebox[\fontcharht\font`X]{A}} +
\colorbox{black}{teal!25}{\makebox[\fontcharht\font`X]{C}} :
Consumer surplus loss \\
\colorbox{black}{violet!10}{\makebox[\fontcharht\font`X]{B}} +
\colorbox{black}{violet!25}{\makebox[\fontcharht\font`X]{F}} :
Producer surplus loss \\
\colorbox{black}{teal!10}{\makebox[\fontcharht\font`X]{A}} +
\colorbox{black}{violet!10}{\makebox[\fontcharht\font`X]{B}} :
Total tax revenue \\
\colorbox{black}{teal!25}{\makebox[\fontcharht\font`X]{C}} +
\colorbox{black}{violet!25}{\makebox[\fontcharht\font`X]{F}} :
Deadweight loss
};

```

Adding this to the graph, we get Figure 6-4.




Figure 6-4: The effect of a gasoline tax on consumption and production and total tax revenue.

And viola! We have added dimension lines and a legend, further communicating what different sections of the graph mean. Now, anyone looking at this graph would pictorially understand that consumers bear the bulk of the tax incidence and the surplus loss from a gasoline excise tax—and why. (Assuming that supply is relatively price-elastic and demand is relatively price-inelastic.) The code that produces Figure 6-4 is:

```
\begin{center}
\hspace*{-3cm}\begin{tikzpicture}
\begin{axis}[
scale = 1.2,
xmin = 0, xmax = 10,
ymin = 0, ymax = 10,
axis lines* = left,
xtick = {0}, ytick = \empty,
clip = false,
]
% Colouring areas
\fill[teal, opacity = 0.1] (0, 3.53) -- (4.14, 3.53) -- (4.14,
5.96) -- (0, 5.96);
\fill[violet, opacity = 0.1] (0, 3.53) -- (4.14, 3.53) -- (4.14,
3) -- (0, 3);
\fill[teal, opacity = 0.25] (4.14, 3.53)-- (5.05, 3.53) -- (4.14,
5.96);
\fill[violet, opacity = 0.25] (4.14, 3)-- (5.05, 3.53) -- (4.14,
3.53);

% Supply and demand curves
\addplot[color = blue, very thick] coordinates {(3,9) (6,1)};
\addplot[color = red, very thick] coordinates {(1,1) (9,6)};
\addplot[color = red, opacity = 0.3, very thick] coordinates
{(1,4) (9,9)};

% Dashed lines
\addplot[color = black, dashed, thick] coordinates {(0, 5.96)
(4.14, 5.96) (4.14, 0)};
\addplot[color = black, dashed, thick] coordinates {(0, 3.53)
(5.05, 3.53) (5.05, 0)};
```

```

\addplot[color = black, dashed, thick] coordinates {(0, 3) (4.1, 3)
};

% Coordinate points
\addplot[color = black, mark = *, only marks, mark size = 3pt]
coordinates {(4.14, 5.96) (5.05, 3.53)};

% Labels
\node [right] at (current axis.right of origin) {Gasoline,  $Q$ };
\node [above] at (current axis.above origin) {Price,  $P$ };
\node [above] at (5.25, 3.6) { $E$ };
\node [above] at (4.35, 6.1) { $E^{\prime}$ };
\node [right] at (6, 1) { $D$ };
\node [right] at (9, 6) { $S$ };
\node [right] at (9, 9) { $S^{\prime}$ };
\node [left] at (0, 3.7) { $P_E$ };
\node [left] at (0, 2.8) { $P_S$ };
\node [left] at (0, 5.96) { $P^{\prime}$ };
\node [below] at (5.05, 0) { $Q_E$ };
\node [below] at (4.14, 0) { $Q^{\prime}$ };
\node [above] at (2, 4.7) { $A$ };
\node [above] at (1, 2) { $B$ };
\node [above] at (4.5, 3.5) { $C$ };
\node [above] at (6.2, 2.2) { $F$ };

% Arrows
\draw[-{Triangle[length=4mm, width=2mm]}, red, opacity = 0.3] (8,
5.8) to (8, 7.8);
\draw[-Triangle] (1.3, 2.4) to [out = 0, in = 270] (2, 3.35);
\draw[-Triangle] (5.9, 2.6) to [out = 180, in = 270] (4.3, 3.4);

% Dimension lines
\draw[|-|] (4.14, -1) to (5.05, -1);
\node [below] at (4.59, -1) {Consumption reduction};
\draw[|-|] (-1, 5.96) to (-1, 3.8);
\node [below, align = left] at (-2.3, 6) {Consumer \\\ incidence};
\draw[|-|] (-1, 3.6) to (-1, 2.8);
\node [below, align = left] at (-2.3, 3.5) {Producer \\\ incidence

```

```

    };
\draw[|-|] (-3.6, 5.96) to (-3.6, 2.8);
\node [below, align = left] at (-4.2, 4.7) {Tax};

% Legend
\node [below right, draw, align = left] at (10.5, 10) {
\colorbox{black}{teal!10}{\makebox[\fontcharht\font`X]{A}} +
\colorbox{black}{teal!25}{\makebox[\fontcharht\font`X]{C}} :
Consumer surplus loss \\
\colorbox{black}{violet!10}{\makebox[\fontcharht\font`X]{B}} +
\colorbox{black}{violet!25}{\makebox[\fontcharht\font`X]{F}}
: Producer surplus loss \\
\colorbox{black}{teal!10}{\makebox[\fontcharht\font`X]{A}} +
\colorbox{black}{violet!10}{\makebox[\fontcharht\font`X]{B}}
: Total tax revenue \\
\colorbox{black}{teal!25}{\makebox[\fontcharht\font`X]{C}} +
\colorbox{black}{violet!25}{\makebox[\fontcharht\font`X]{F}}
: Deadweight loss
};
\end{axis}
\end{tikzpicture}\hspace*{-3cm}
\end{center}

```

7 Conclusion

With this toolbox, you can create virtually any economic graph. The effort to learn pgfplots and LaTeX is worth it. While on one hand, most word processors such as Google Docs (with Google Sheets) and Word (with Excel) can provide the necessities of economic graphs—axes, lines, curves, and labels—pgfplots has more flexibility, from shading areas, to adding dashed lines, and to adding dimension lines. All done in style.

Of course, since it is LaTeX, this is only the tip of the iceberg. Everyone stands on the shoulders of giants. First, David Knuth created TeX to typeset documents. Then Leslie Lamport created LaTeX as a markup to TeX. Then Till Tantau created TikZ on top of this foundation to create graphics. After, before, and between this

sequence is a list of many other contributors to the giant on which we stand. Finally, we come along and, with serendipity, find out that we can create beautiful and sleek graphs for economics.

Yet, we can always do more. For instance, we could automate the graph creation by defining variables and automatically calculating the points of intersection. That, however, is outside the scope of this guide.

If I have done my job successfully, you have partially scaled LaTeX's reputable learning curve in this one domain of making graphs. If I have not, you have been left abandoned, with no guidance on how to begin your ascent (at least, not from me!). Let us hope for the former.

8 Resources

Getting started with LaTeX and learning

- TeX Live: <http://www.tug.org/texlive/>.
This is a portal to install TeX, the TeXworks IDE, and a few ubiquitous packages (including pgfplots and xcolor).
- Overleaf: <https://www.overleaf.com/>.
If installing LaTeX is too much of a commitment, there are many online LaTeX editors that are sufficient for this guide. Overleaf has pgfplots and xcolor already installed, so you do not need to worry about that. On the subject of pricing, they offer a free version which contains all the essentials and only lacks a few specialized features such as GitHub integration or real-time collaboration.
- L^AT_EX on Wikibooks: <https://en.wikibooks.org/wiki/LaTeX>.
This guide explains the basics of how LaTeX works. From installing LaTeX, to making a document, to more technical aspects such as boxes, the Wikibook has its explanations and accompanying code.
- TeX-LaTeX Stack Exchange: <https://tex.stackexchange.com/>.
Stack Exchange is a forum where people can ask and answer questions. This

subforum specifically deals with LaTeX, which includes pgfplots. Of course, if you have a question, first use the search feature to check if it has been answered before. Then, read the rules and ask away.

Pgfplots and TikZ guides

- Pgfplots package guide by Overleaf: https://www.overleaf.com/learn/latex/Pgfplots_package. Pgfplots was designed to make line graphs and scatter plots. We had to adapt it for economic graphs. But if you want to use pgfplots for more conventional purposes, read this guide by Overleaf.
- Usepackage{TikZ} for economists by Kevin Goulding:
<http://static.latexstudio.net/wp-content/uploads/2016/06/tikzforeconomists-110619150244-phpapp01.pdf>. This is another guide on making economic graphs in LaTeX. It uses TikZ, however, and was written before pgfplots was released. Since pgfplots is a dependency of TikZ, the code explained in Kevin Goulding's guide may be applicable to making economic graphs with pgfplots.

Package installation links and documentation

Although the LaTeX editing solutions above—TeX Live and Overleaf—already have these packages pre-installed, if you do not have these packages or want the latest versions, you can access them at these CTAN links. Also, package manuals are available on CTAN, but beware they are hundreds of pages of technical documentation.

- Pgfplots package on CTAN: <https://ctan.org/pkg/pgfplots>. This pgfplots package installation link also comes with TikZ.
- Xcolor package on CTAN: <https://ctan.org/pkg/xcolor>.

9 Other examples

Again, all of the finished graphs in this guide—the PDFs and the code—are available on the GitHub repository, including the examples below. The repository can be accessed at <https://github.com/jackypacky/pgf-econ-graphs>.




Figure 9-1: The effect of rent control on the rental market.




Figure 9-2: The effect of tariffs in a small economy on the domestic market (L) and the import market (R).




Figure 9-3: The market equilibrium, W' , in an Edgeworth box between two consumers, A and B , and two commodities, x and y .